YOGANANDA GOVIND

AI & Cloud Engineer | AWS Solutions Architect | Serverless & AI Platform Builder
UK | Available for Contract or Permanent | Remote | Hybrid

Phone: +44 07365558999
Email: yoganand.govind@gmail.com
Linkedin: www.linkedin.com/in/yoganandgovind
Portfolio: https://www.yoganandgovind.me

SUMMARY

Enterprise software engineer transitioning into AI and cloud engineering, with 12+ years of experience building large-scale backend systems and cloud-native and enterprise applications. Proven experience designing and implementing end-to-end architectures, including a multi-tenant AI-powered SaaS platform on AWS.

Strong background in serverless, event-driven systems, and microservices, with hands-on expertise in defining High-Level Design (HLD) and Low-Level Design (LLD), system trade-offs, and scalable data architectures. Experienced in integrating AI/ML services and modernising legacy systems using cloud-native patterns.

CORE COMPETENCIES

Cloud & Serverless Architecture

- AWS (Lambda, API Gateway, DynamoDB, S3, Step Functions, SQS)
- Event-driven systems & async processing
- Multi-tenant SaaS architecture
- Infrastructure as Code (AWS CDK), Terraform (working knowledge)
- Cost optimisation & observability

AI & Intelligent Automation

- Amazon Bedrock & Textract
- LLM-powered document intelligence
- AI workflow orchestration
- Scalable document processing pipelines

Backend & Distributed Systems

- Java, TypeScript, Node.js
- RESTful services & system integrations
- High-volume data processing
- Performance tuning & reliability engineering

DevOps & Delivery

- CI/CD automation
- Production support in SLA-driven environments
- Agile / Scrum delivery

WORK EXPERIENCE

Senior Software Engineer

Mphasis UK Limited, London, UK
2014-Present

Project - Personal Independence Payment (UK public sector)

- Delivered end-to-end **Cúram SPM solutions** across multiple modules, supporting millions of PIP claims annually
- Led development across **Case Management, Financials, Task Management, and Batch Framework**, ensuring scalable and maintainable solutions
- Translated complex business and policy requirements into rule-driven system implementations (**CER**)
- Optimised **batch processing and SQL queries**, improving performance for high-volume workloads (millions of records)
- Resolved **Sev1/Sev2 production incidents** within strict SLA timelines in a high-availability environment

- Contributed to **production deployments, go-live support, and post-release validation,** ensuring minimal disruption to live services
- Collaborated with BAs, architects, and product owners to deliver iterative enhancements across multiple release cycles
- Established development best practices and conducted code reviews, improving code quality and maintainability
- Mentored junior developers and supported onboarding, improving team delivery capability
- Developed integrations with external systems using **REST/SOAP** services and custom Java components
- Managed CI/CD pipelines using **Jenkins, GitLab, TortoiseGit and SVN,** ensuring consistent and reliable deployments
- Conducted **R&D and proof-of-concept** development to evaluate new features and implementation approaches

Featured Project

Autowired.ai – AI-Powered Document Intelligence Platform

- Created end-to-end **HLD** and **LLD** for a multi-tenant serverless SaaS platform on AWS
- Defined **system architecture** including API Gateway, Lambda orchestration, event-driven pipelines, and DynamoDB data models
- Built an **event-driven processing pipeline** using SNS and SQS for scalable workflows
- Implemented **AI-powered extraction** using **Amazon Textract** and **Bedrock**
- Evaluated design trade-offs between **synchronous vs asynchronous processing**, optimising for **scalability**, **cost**, and **reliability**
- Designed **multi-tenant architecture with secure data isolation**
- Implemented **stateless compute** and **horizontal scaling patterns**

Key Architecture Highlights:

Multi-Tenant SaaS Architecture

- Implemented strict tenant isolation using DynamoDB single-table design with partition key scoping and tenant-aware access control. Ensured secure data separation and scalable onboarding.

Asynchronous Batch Processing Pipeline

- Built an event-driven document workflow using S3 triggers, SQS queues, AWS Lambda, and Step Functions to process high document volumes reliably without blocking user interactions.

Token-Aware AI Processing

- Integrated Amazon Bedrock with controlled prompt orchestration and intelligent caching (24hr TTL) to reduce repeated inference calls and optimise token consumption, lowering AI costs by ~40%.

Subscription & Usage-Based Model

- Integrated Stripe for subscription management, usage limits, and tier-based document processing quotas. Enforced plan restrictions at API and processing layers.

Cost-Optimised Serverless Design

- Fully serverless stack (Lambda, DynamoDB, S3, API Gateway) enabling automatic scaling with minimal idle cost.

Live: https://autowired.ai

CERTIFICATIONS

- AWS Certified Generative AI Developer – Professional (2026–2029)
- AWS Certified Solutions Architect – Associate (2025-2028)
- AWS Certified Developer – Associate (2021-2024)
- IBM Cúram V6 Application Developer (IBM certified)

EDUCATION

- Master of Computer Application (MCA) - 2014
- Bachelor of Computer Application (BCA) - 2011

ADDITIONAL INFORMATION

- Extensive experience delivering UK public sector, policy-driven benefit systems in high-availability environments
- Experience transitioning from legacy enterprise systems to cloud-native architectures
- Strong understanding of security-conscious, SLA-driven environments
- Experience building both enterprise platforms and startup SaaS products
- Containers & Orchestration: Docker, Kubernetes (fundamentals, currently upskilling)
- Comfortable in remote, distributed teams