

YOGANANDA GOVIND

AI & Cloud Engineer | AWS Solutions Architect | Serverless & AI Platform Builder
UK | Available for Contract or Permanent | Remote | Hybrid

Phone: +44 07365558999

Email: yoganand.govind@gmail.com

Linkedin: www.linkedin.com/in/yoganandgovind

Portfolio: <https://www.yoganandgovind.me>

SUMMARY

Enterprise software engineer transitioning into AI and cloud engineering, with 12+ years of experience building large-scale backend systems and cloud-native and enterprise applications. Proven experience designing and implementing end-to-end architectures, including a multi-tenant AI-powered SaaS platform on AWS.

Strong background in serverless, event-driven systems, and microservices, with hands-on expertise in defining High-Level Design (HLD) and Low-Level Design (LLD), system trade-offs, and scalable data architectures. Experienced in integrating AI/ML services and modernising legacy systems using cloud-native patterns.

CORE COMPETENCIES

Cloud & Serverless Architecture

- AWS (Lambda, API Gateway, DynamoDB, S3, Step Functions, SQS)
- Event-driven systems & async processing
- Multi-tenant SaaS architecture
- Infrastructure as Code (AWS CDK), Terraform (working knowledge)
- Cost optimisation & observability

AI & Intelligent Automation

- Amazon Bedrock & Textract
- LLM-powered document intelligence
- AI workflow orchestration
- Scalable document processing pipelines

Backend & Distributed Systems

- Java, TypeScript, Node.js
- RESTful services & system integrations
- High-volume data processing
- Performance tuning & reliability engineering

DevOps & Delivery

- CI/CD automation
- Production support in SLA-driven environments
- Agile / Scrum delivery

WORK EXPERIENCE

Senior Software Engineer | Team Lead

Mphasis UK Limited, London, UK | 2014-Present

Project – Personal Independence Payment (UK public sector)

- Delivered end-to-end Cúram SPM solutions for major DWP benefit programmes, including **Universal Credit and Personal Independence Payment (PIP)**, supporting millions of claims annually in a **large-scale government environment**
- Led backend development using **Java, Spring Boot, REST APIs, and microservices-based architectures** to build scalable and maintainable enterprise applications, while currently working within a **user research-focused team managing 3 team members**
- Developed and enhanced responsive frontend applications using **React and Next.js**, improving user experience and operational efficiency
- Worked across **Case Management, Financials, Task Management, and Batch Framework**, ensuring robust system design and high availability
- Translated complex business and policy requirements into rule-driven system implementations using CER (Cúram Express Rules)
- Designed and implemented integrations with external systems using **REST/SOAP** services and custom Java components
- Optimised batch processing, SQL queries, and backend services, improving performance for high-volume workloads handling millions of records

- Resolved **Sev1/Sev2** production incidents within strict SLA timelines in a mission-critical, high-availability environment
- Contributed to **production deployments, go-live support, release validation, and environment stability** with minimal disruption to live services
- Managed and supported **CI/CD pipelines using Jenkins, GitLab, SVN, and TortoiseGit**, enabling reliable and automated deployments
- Collaborated with **business analysts, architects, product owners, end-users, quality analysts and cross-functional teams** to deliver iterative enhancements across multiple release cycles
- Currently lead a team of 3 members within a **user research and development-focused team, driving the delivery of MVPs** and new feature implementations by collaborating closely with stakeholders to understand requirements, gather feedback, and build scalable solutions
- Led **R&D initiatives and proof-of-concept development** to evaluate modernisation approaches, deliver new platform capabilities, and accelerate feature delivery through scalable solution designs
- Established development best practices, conducted code reviews, and improved application maintainability and engineering standards
- Mentored junior developers and supported onboarding activities, improving overall team capability and delivery efficiency
- Worked within **Agile/Scrum** delivery teams, contributing to sprint planning, estimation, technical design discussions, and continuous improvement initiatives

Featured Project

Autowired.ai – AI-Powered Document Intelligence Platform

- Created end-to-end **HLD** and **LLD** for a multi-tenant serverless SaaS platform on AWS
- Defined **system architecture** including API Gateway, Lambda orchestration, event-driven pipelines, and DynamoDB data models
- Designed and provisioned AWS infrastructure for Autowired.ai using **AWS CDK**, later migrating and managing the infrastructure with **Terraform** to strengthen Infrastructure-as-Code practices and cloud automation expertise
- Built an **event-driven processing pipeline** using SNS and SQS for scalable workflows
- Implemented **AI-powered extraction** using **Amazon Textract** and **Bedrock**
- Evaluated design trade-offs between **synchronous vs asynchronous processing**, optimising for **scalability, cost, and reliability**
- Designed **multi-tenant architecture with secure data isolation**
- Implemented **stateless compute** and **horizontal scaling patterns**
- Built a production **RAG-powered chat system** enabling customers to query structured business data in natural language, reducing support overhead by surfacing contextual answers directly from **indexed knowledge bases using vector similarity search**.
- Developed a **context-aware conversational AI feature** for a marketing platform using Retrieval-Augmented Generation (RAG), grounding LLM responses in proprietary marketing content and campaign data to improve response accuracy and brand consistency.

Key Architecture Highlights:

Multi-Tenant SaaS Architecture

- Implemented strict tenant isolation using DynamoDB single-table design with partition key scoping and tenant-aware access control. Ensured secure data separation and scalable onboarding.

Asynchronous Batch Processing Pipeline

- Built an event-driven document workflow using S3 triggers, SQS queues, AWS Lambda, and Step Functions to process high document volumes reliably without blocking user interactions.

Token-Aware AI Processing

- Integrated Amazon Bedrock with controlled prompt orchestration and intelligent caching (24hr TTL) to reduce repeated inference calls and optimise token consumption, lowering AI costs by ~40%.

Subscription & Usage-Based Model

- Integrated Stripe for subscription management, usage limits, and tier-based document processing quotas. Enforced plan restrictions at API and processing layers.

Cost-Optimised Serverless Design

- Fully serverless stack (Lambda, DynamoDB, S3, API Gateway) enabling automatic scaling with minimal idle cost.

Live: <https://autowired.ai>

CERTIFICATIONS

- AWS Certified Generative AI Developer – Professional (2026–2029)
- AWS Certified Solutions Architect – Associate (2025–2028)
- AWS Certified Developer – Associate (2021–2024)
- IBM Cúram V6 Application Developer (IBM certified)

EDUCATION

- Master of Computer Application (MCA) - 2014
- Bachelor of Computer Application (BCA) - 2011

ADDITIONAL INFORMATION

- Extensive experience delivering UK public sector, policy-driven benefit systems in high-availability environments
- Experience transitioning from legacy enterprise systems to cloud-native architectures
- Strong understanding of security-conscious, SLA-driven environments
- Experience building both enterprise platforms and startup SaaS products
- Containers & Orchestration: Docker, Kubernetes (fundamentals, currently upskilling)
- Comfortable in remote, distributed teams